


# Querying Dynamic Tables with SQL

---

Flink SQL Training

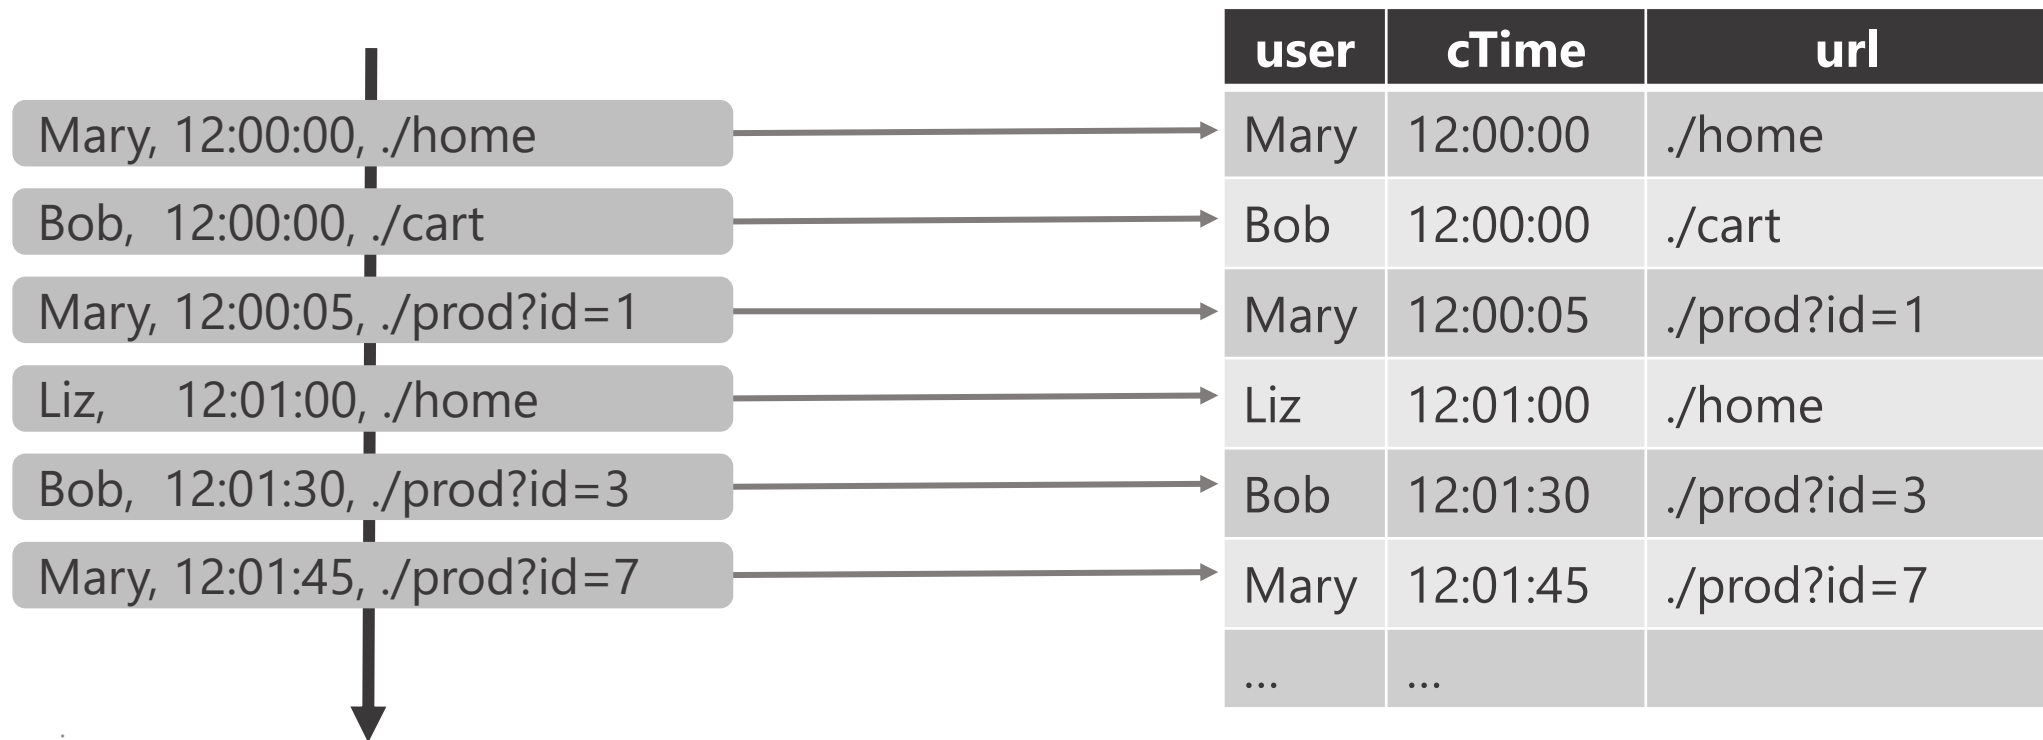
<https://github.com/ververica/sql-training>

# Continuous Queries in Flink

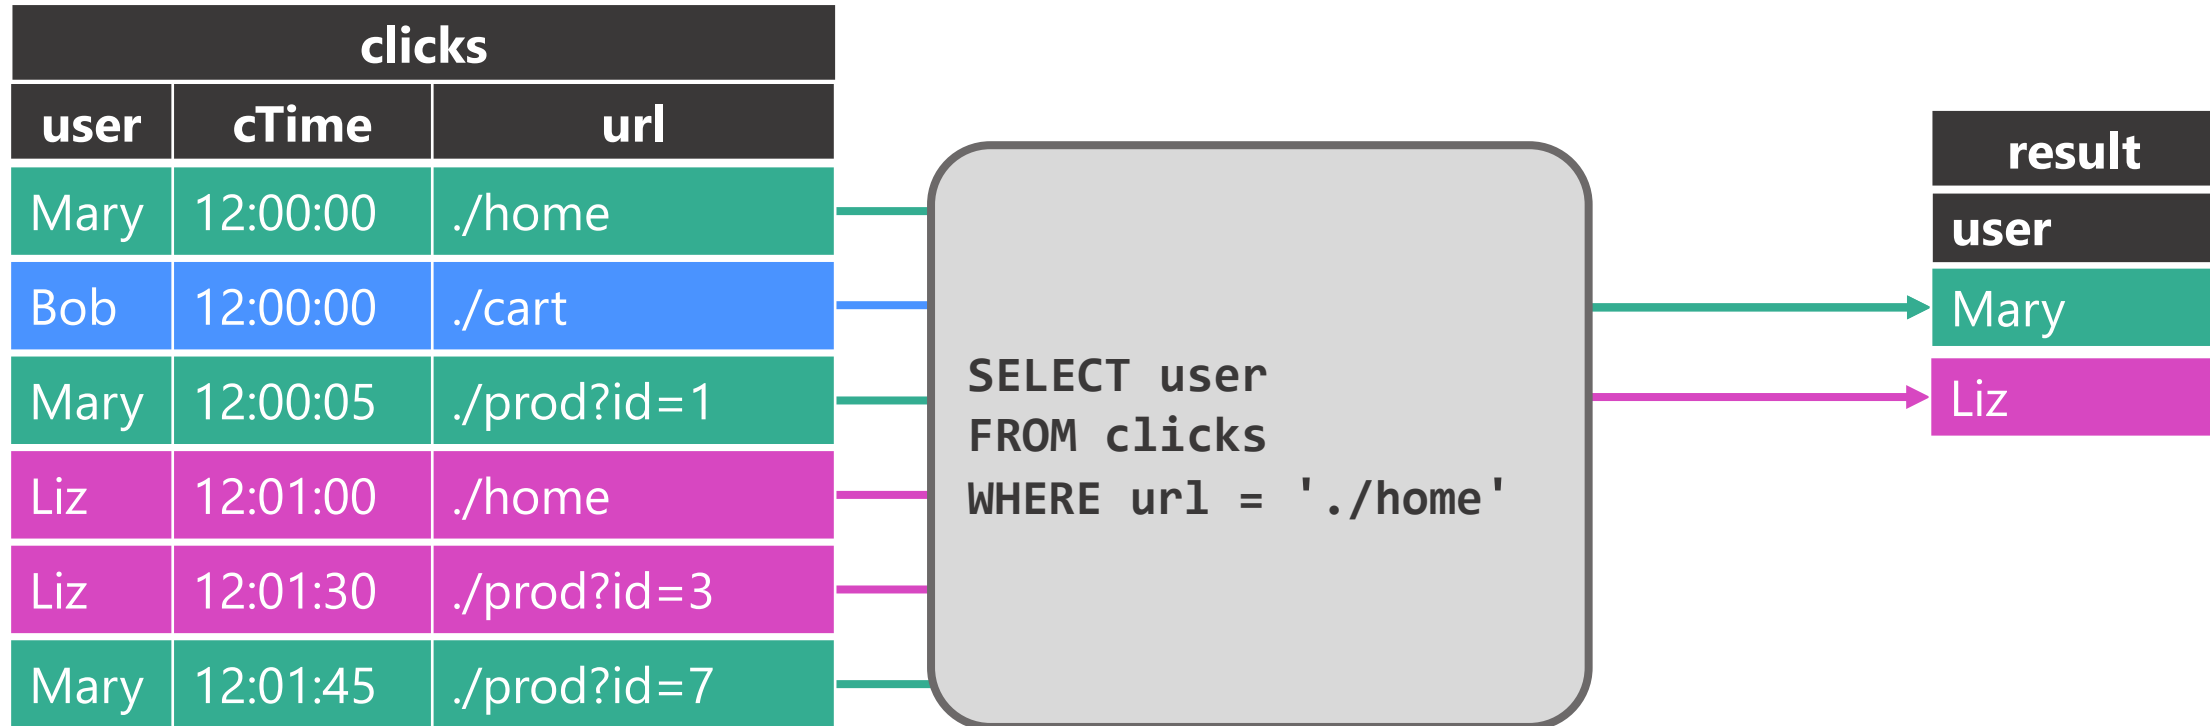


# Stream → Dynamic Table: INSERT

- Append mode
  - Stream records are appended to table
  - Table grows as more data arrives

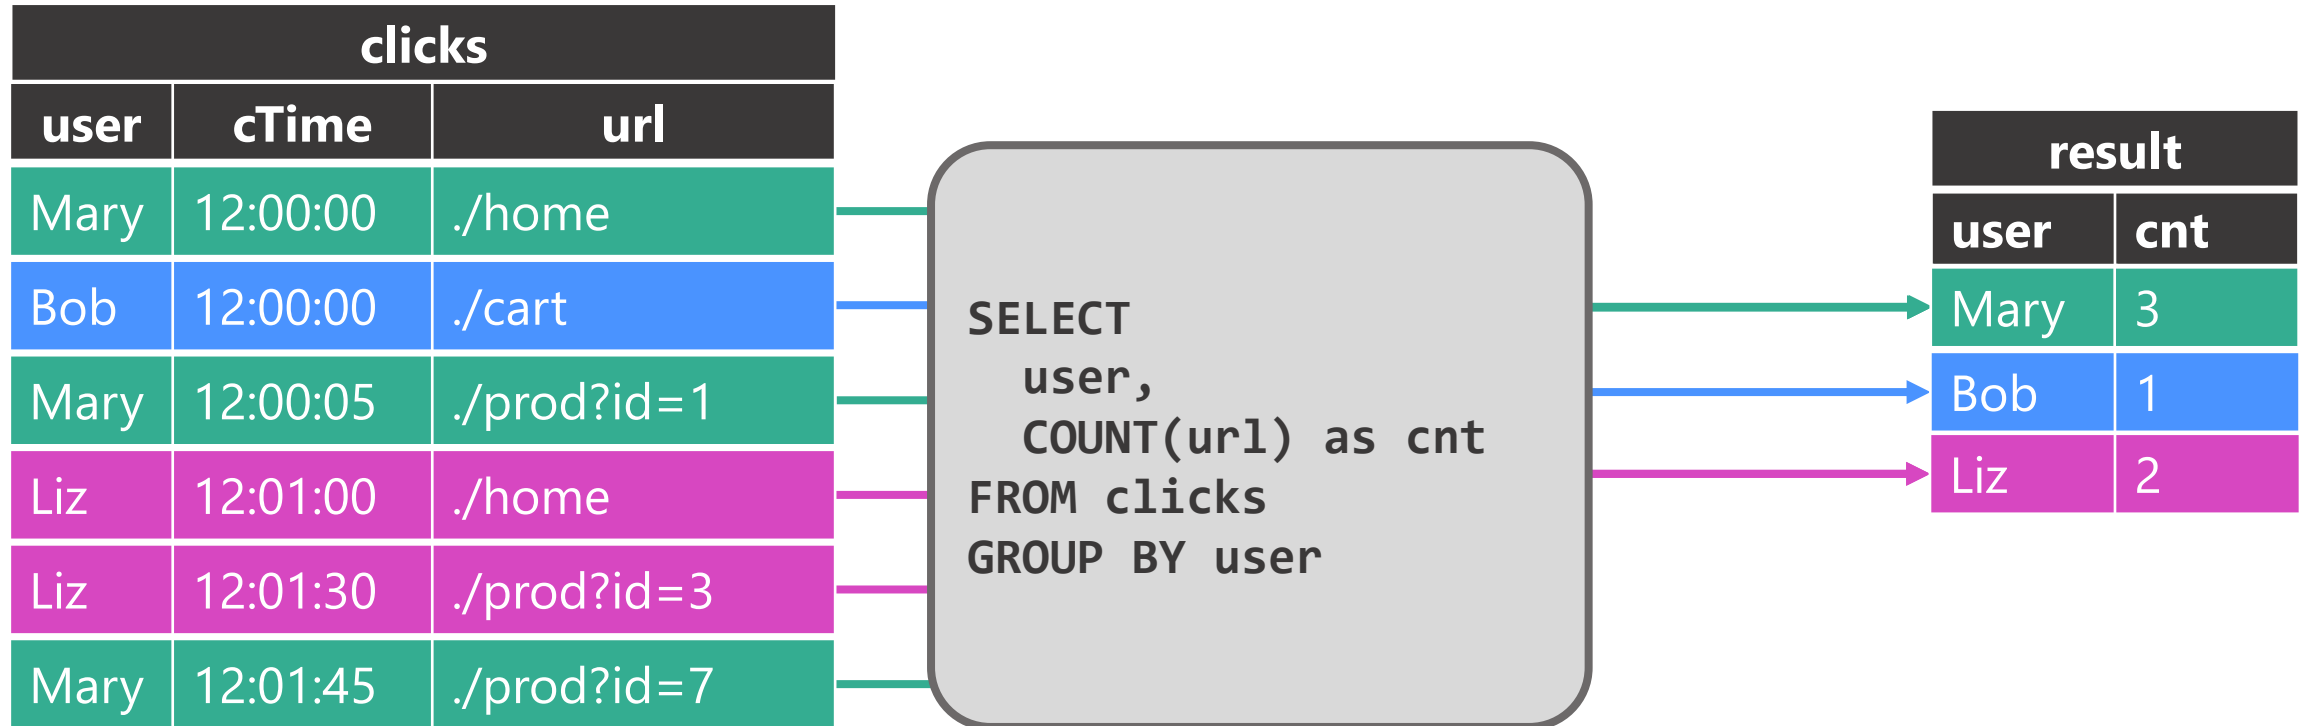


# Querying a Dynamic Table



Rows of result table are appended.

# Querying a Dynamic Table



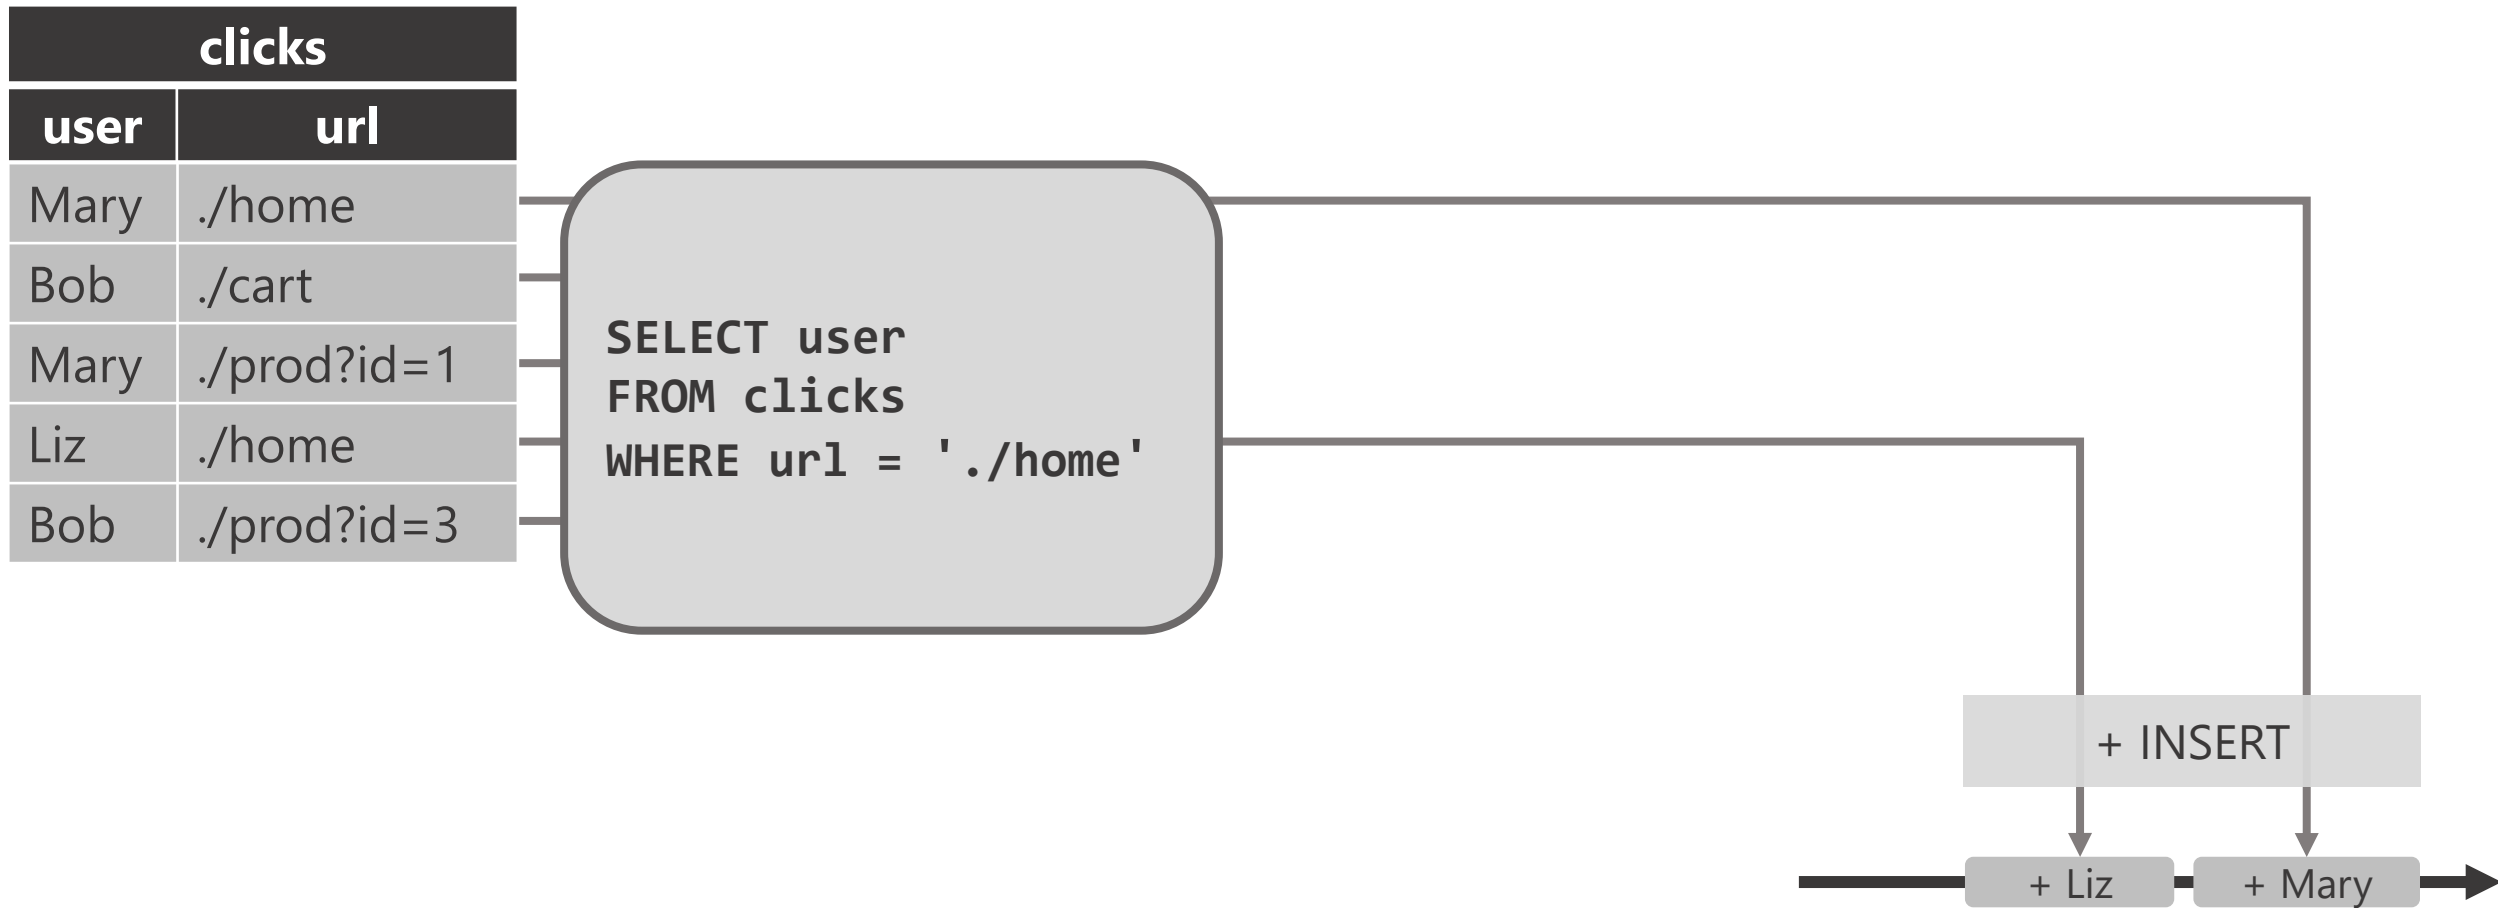
Rows of result table are updated.

# Dynamic Table → Stream

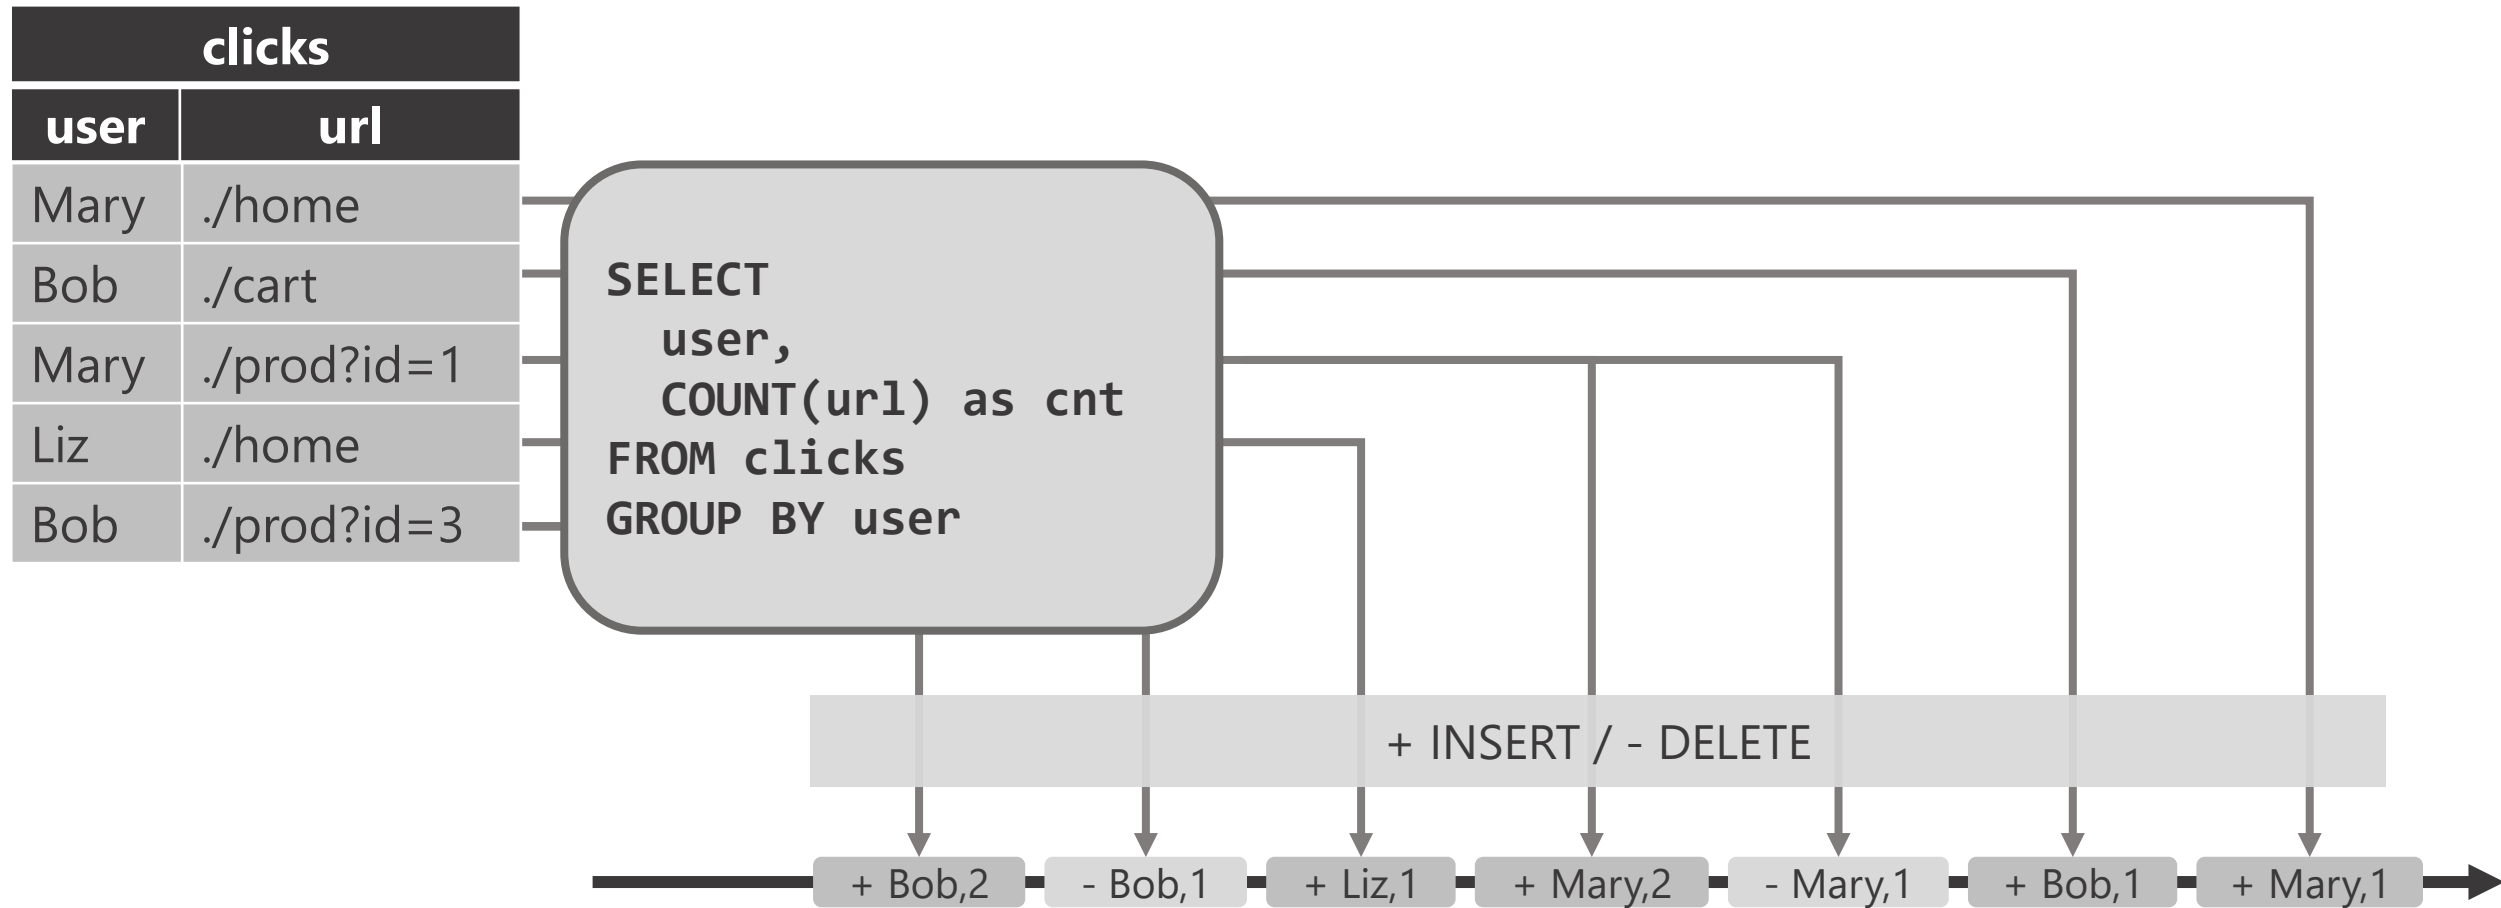
- Converting a dynamic table into a stream
- Dynamic tables might update or delete existing rows
- Updates must be encoded in outgoing stream



# Dynamic Table → Stream: INSERT-only

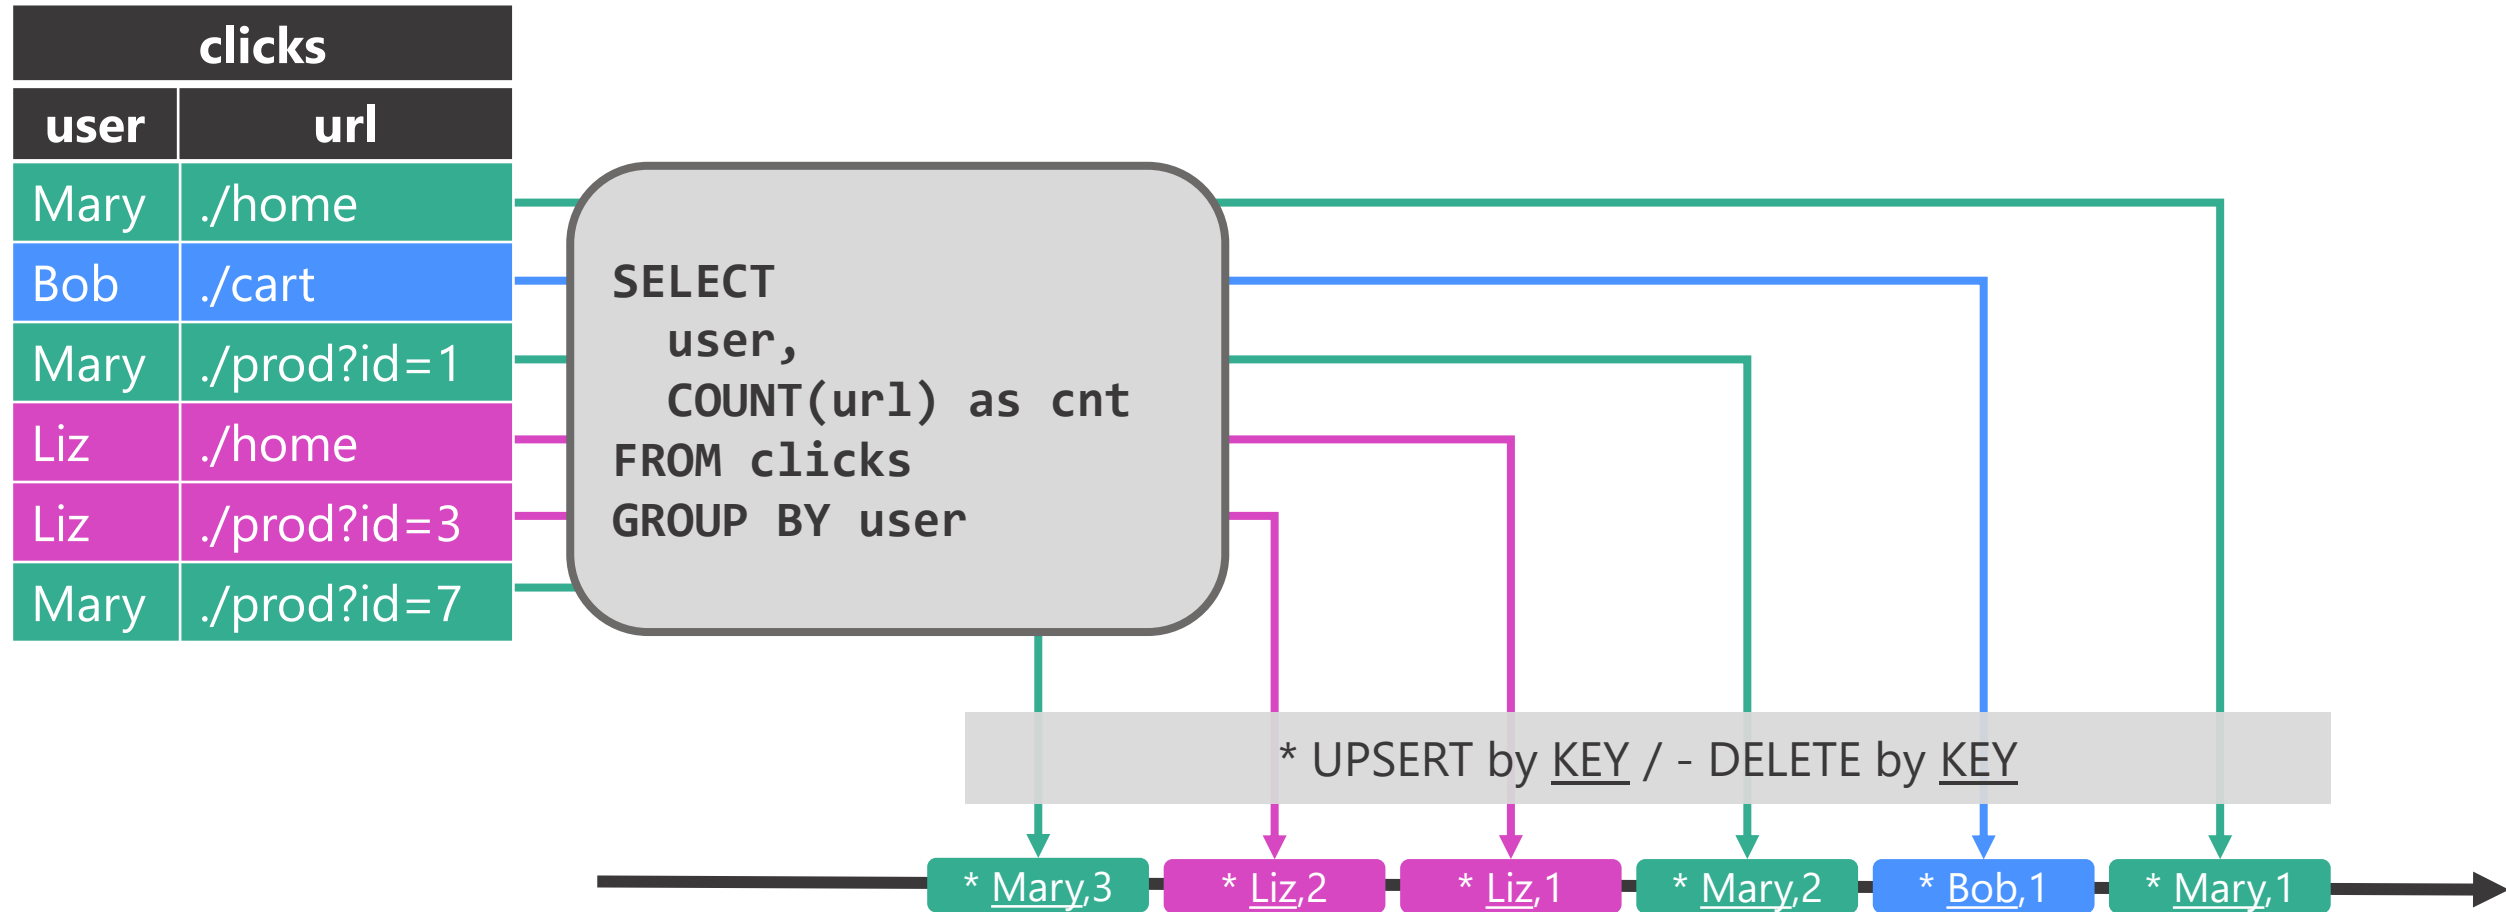


# Dynamic Table → Stream: INSERT+DELETE





# Dynamic Table → Stream: UPSERT+DELETE



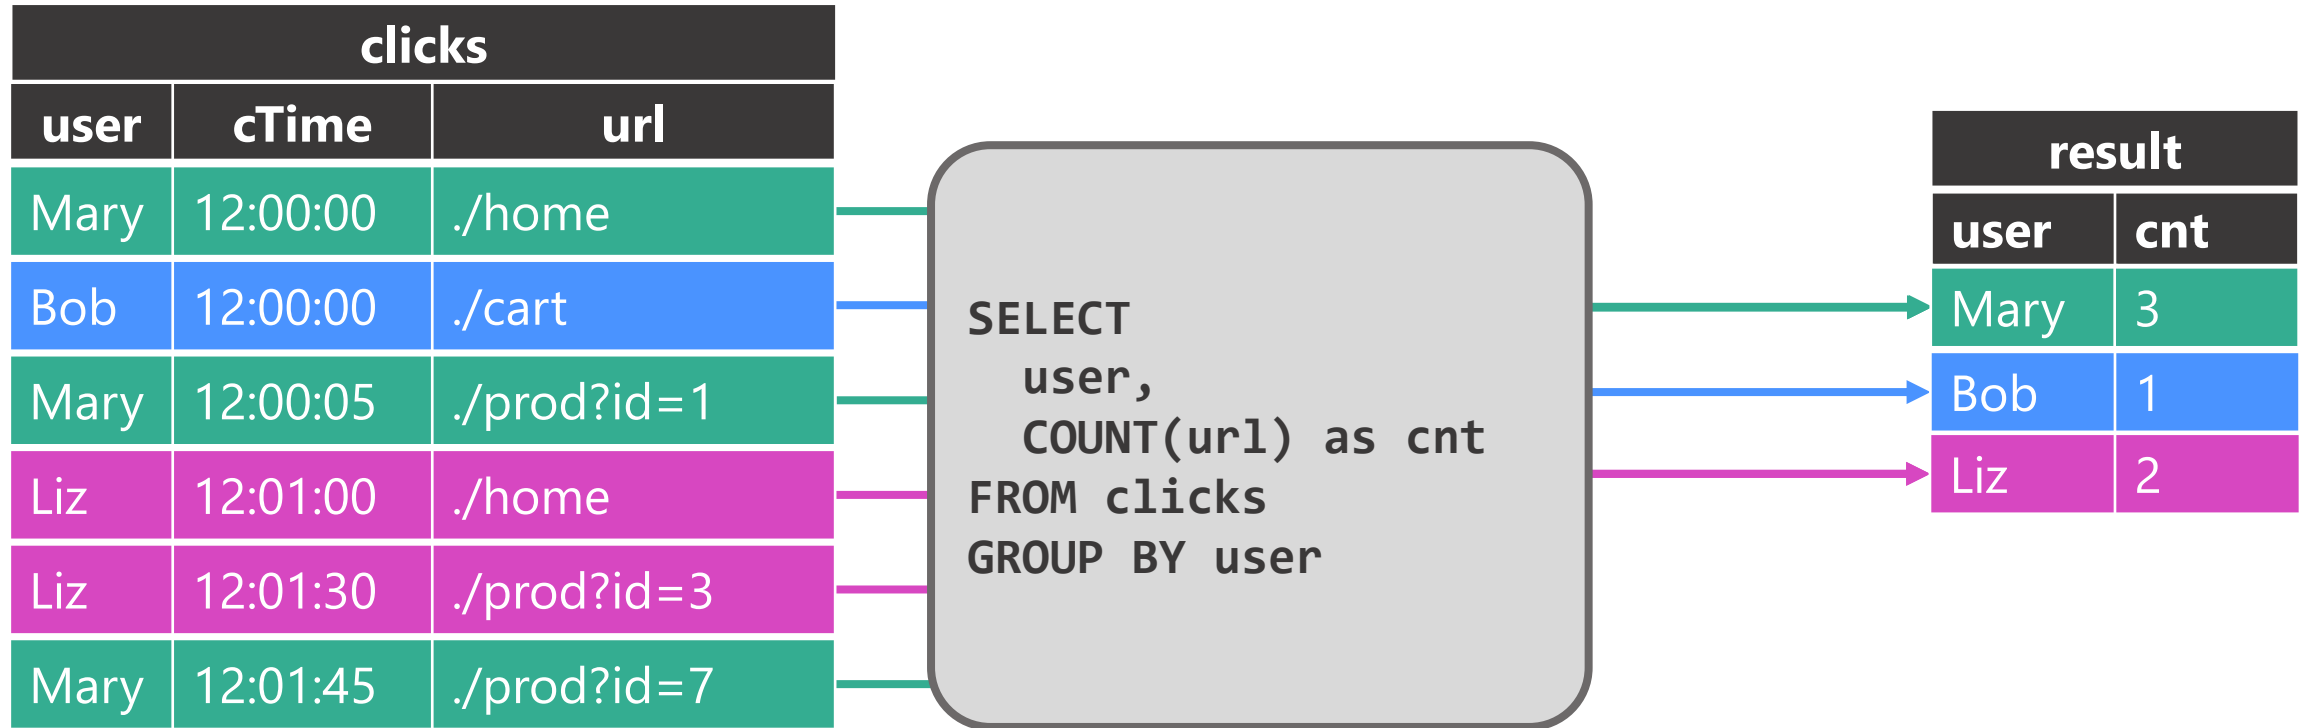
# Operators and State

# Operator Types

- Stateless operators
  - Filter
  - Projection
- Materializing operators
  - Aggregation
  - Joins
- Temporal operators (discussed in later sessions)
  - Window aggregation (GROUP BY, OVER)
  - Time-based joins (Interval join, Temporal-table join)
  - Pattern matching (MATCH\_RECOGNIZE)



# Materializing Aggregation



Rows of result table are updated.

# Materializing Aggregation

```
SELECT
    user,
    COUNT(url) as cnt
FROM clicks
GROUP BY user
```

- The aggregation needs to maintain a count for every user forever.
  - Every user could click at any point in time
- The aggregation state is growing with every new user
  - For some aggregation functions, state grows with every new input row



# Stateful Operators

- Materializing operators
  - Computations are not bound by temporal condition and never complete
  - Input and output records can be updated or deleted
  - Hold records and/or results forever in state
  - State can grow over time (depending on query and data)
- Temporal operators
  - Associate records based on a temporal condition
  - Only accept new records. Previously added records can not be updated or deleted.
  - Hold records and/or results in state until a computation is complete
  - Automatically clean up state as soon as records and results are no longer needed



# Managing State Size of Materializing Operators

- Query state might grow indefinitely
  - Depends on query and input tables
- Slowly growing state can be addressed by scaling the query
  - `SELECT user, COUNT(*) FROM logins GROUP BY user;`
- State can be automatically pruned
  - `SELECT session, COUNT(*) FROM clicks GROUP BY session;`
  - Rows and persisted results can be removed after an idle timeout



# Idle State Clean Up

- Configure Flink to automatically remove state that was not accessed for x time
  - The query result is not updated when state is removed
- Query result remains consistent if removed state is not needed again
- Query result becomes inconsistent if query needs to access state that was removed!
- Trade the accuracy of the result for size of state





# Summary

# Summary

- Streams are interpreted as changelog for a Table
  - Flink 1.10 supports INSERT-only stream to table conversion
  - Flink 1.11 will support full changelog conversion
- SQL queries on dynamic tables yield another dynamic table
  - Input and query determines whether resulting dynamic table is append-only or updating
- Dynamic tables can be converted back into streams
  - INSERT-only, INSERT+DELETE, UPSERT+DELETE



# Summary

- Use regular SQL to run queries on dynamic tables
  - No need to learn special syntax or semantics
  - Writing and executing queries is easy, BUT...
- Pay attention to the state requirements of your query
  - Depending on query and input data, the state might grow indefinitely
  - Enable idle state pruning to trade off query accuracy and state consumption



# Hands On Exercises

# Querying Dynamic Tables with SQL

Continue with the hands-on exercises in  
"Querying Dynamic Tables with SQL"

<https://github.com/ververica/sql-training/wiki/Querying-Dynamic-Tables-with-SQL>

We are here to help!





# ververica

---

[www.ververica.com](http://www.ververica.com)

@VervericaData